Autonomy Allocation Matrix v1

Normative Requirements (v1.0)

• Autonomy is allocated per decision, not per system.

• Five autonomy levels range from human-executed to AI-executed and AI-monitored.

• Higher autonomy requires stronger evidence, not higher confidence.

• Autonomy must regress as easily as it advances.

• Human accountability never transfers to machines.

Interpretive Guidance (Non-Normative)

The Autonomy Allocation Matrix replaces vague human-in-the-loop assumptions with explicit authority boundaries.

Most AI failures stem from premature autonomy, not model capability.

Regression is a safety mechanism, not a failure state.

Autonomy is a risk allocation decision requiring evidence.